

# **CS 190I**

# **Deep Learning**

# **Regularization**

Lei Li (leili@cs)  
UCSB

Acknowledgement: Slides borrowed from Bhiksha Raj's 11485 and Mu Li & Alex Smola's 157 courses on Deep Learning, with modification


# Recap

---

- Back propagation for Feed-forward neural network
- Model evaluation
- Cross validation
- Overfitting and underfitting

# Underfitting and Overfitting


---



Underfitting



Desired




Overfitting

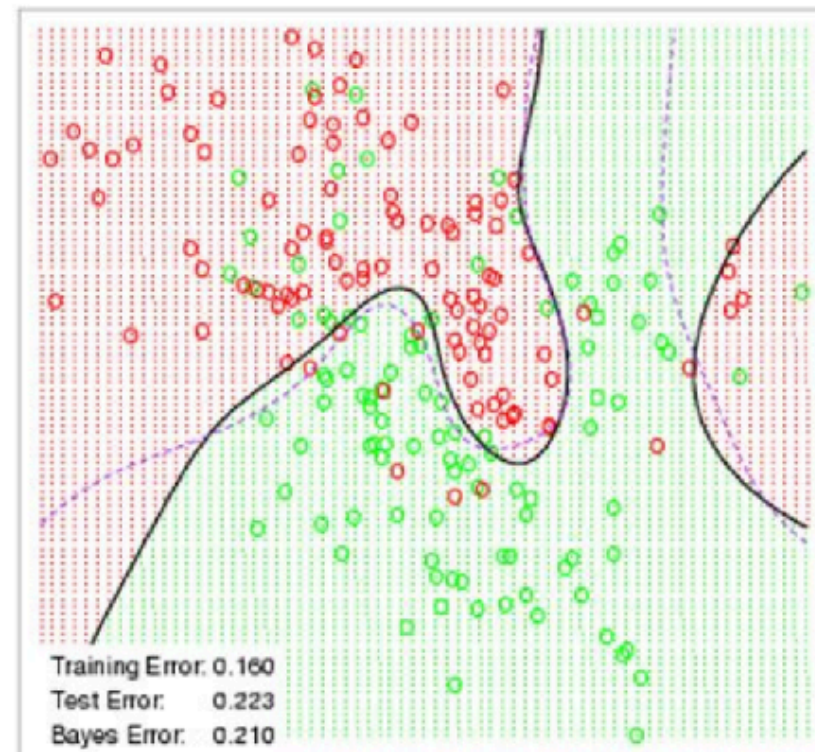
Image credit: [hackernoon.com](https://hackernoon.com)

# Regularization

Neural Network - 10 Units, No Weight Decay



Neural Network - 10 Units, Weight Decay=0.02




# L<sub>2</sub> Regularization as Hard Constraint

- Reduce model complexity by limiting value range

$$\min \ell(\theta) \quad \text{subject to} \quad \|\theta\|^2 \leq \lambda$$

- Often do not regularize bias  $b$ 
  - Doing or not doing has little difference in practice
- A small  $\lambda$  means more regularization



# L<sub>2</sub> Regularization as Soft Constraint

---


- Using Lagrangian multiplier method
- Minimizing the loss plus additional penalty

$$\min \ell(\theta) + \frac{\lambda}{2} \|\theta\|^2$$

- Hyper-parameter  $\lambda$  controls regularization importance
- $\lambda = 0$ : no effect
- $\lambda \rightarrow \infty, \theta^* \rightarrow \mathbf{0}$

# Illustrate the Effect on Optimal Solutions

$$\mathbf{w}^* = \arg \min \frac{1}{2} \|w^T x + b - y\|_2^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$



# Update Rule - Weight Decay

---

- Compute the gradient

$$\frac{\partial}{\partial \theta} \left( \ell(\theta) + \frac{\lambda}{2} \|\theta\|^2 \right) = \frac{\partial \ell(\theta)}{\partial \theta} + \lambda \theta$$

- Update weight at step  $t$

$$\theta_{t+1} = (1 - \eta\lambda)\theta_t - \eta \frac{\partial \ell(\theta_t)}{\partial \theta_t}$$

backprop



- Often  $\eta\lambda < 1$ , so also called weight decay in deep learning



# Weight Decay in Pytorch

---

```
import torch
```

```
learning_rate = 1e-3
```

```
weight_decay = 1.0
```

```
optimizer =
```

```
torch.optim.SGD(model.parameters())
```

```
, lr=learning_rate,
```

```
weight_decay=weight_decay)
```

# General Penalty

---

- Minimizing the loss plus additional penalty

$$\min \mathcal{L}(\theta) + R(\theta)$$

- $\mathcal{L}(\theta)$  is the original loss
- $R(\theta)$  is penalty (or regularization term), not necessary smooth

# L1 Regularization

---

- Minimizing the loss plus additional penalty

$$\min \ell(\theta) + \lambda |\theta|$$

- $\ell(\theta)$  is the original loss
- using L1 norm as penalty

# L1 Update Rule - Soft Thresholding

---

- $\ell(\theta) + \lambda |\theta|$  is not always differentiable!
- Soft-threshold (Proximal operator):

$$S_\lambda(x) = \text{sign}(x) \max(0, |x| - \lambda) = \text{sign}(x) \text{Relu}(|x| - \lambda)$$

- Update weight at step  $t$

$$\tilde{\theta}_t = \theta_t - \eta \frac{\partial \ell(\theta_t)}{\partial \theta_t}$$

$$\theta_{t+1} = S_\lambda(\tilde{\theta})$$

- Also known as Proximal Gradient Descent

# Effects of L1 and L2 Regularization

---

- L1 Regularization
  - will make parameters sparse (many parameters will be zeros)
  - could be useful for model pruning
- L2 Regularization
  - will make the parameter shrink towards 0, but not necessary 0.

# Dropout



# Motivation

---

- A good model should be robust under modest changes in the input
  - Dropout: inject noises into internal layers (simulating the noise)



# Add Noise without Bias

---

- Add noise into  $\mathbf{x}$  to get  $\mathbf{x}'$ , we hope

$$\mathbf{E}[\mathbf{x}'] = \mathbf{x}$$

- Dropout perturbs each element by

$$x'_i = \begin{cases} 0 & \text{with probability } p \\ \frac{x_i}{1-p} & \text{otherwise} \end{cases}$$



# Apply Dropout


- Often apply dropout on the output of hidden fully-connected layers

$$\mathbf{h} = \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

$$\mathbf{h}' = \text{dropout}(\mathbf{h})$$

$$\mathbf{o} = \mathbf{W}_2 \mathbf{h}' + \mathbf{b}_2$$

$$\mathbf{y} = \text{softmax}(\mathbf{o})$$



# Dropout in Training and Inference

---

- Dropout is only used in training

$$\mathbf{h}' = \text{dropout}(\mathbf{h})$$


- No dropout is applied during inference!
- Pytorch Layer:

```
torch.nn.Dropout(p=0.5)
```

# Dropout: Typical results

- From Sr  
architect  
– 2-4 hidc

r different  
ropout



# Recap

---

- Regularization
  - to avoid model overfitting
  - L1 ==> more sparse parameters
  - L2/Weight decay ==> shrink parameters
  - Dropout, equivalent to L2, but as a network Layer

# Numerical Stability

# Gradients for Neural Networks

---

- Consider a network with  $d$  layers

$$\mathbf{h}^t = f_t(\mathbf{h}^{t-1}) \quad \text{and} \quad y = \ell \circ f_d \circ \dots \circ f_1(\mathbf{x})$$

- Compute the gradient of the loss  $\ell$  w.r.t.  $\mathbf{W}_t$

$$\frac{\partial \ell}{\partial \mathbf{W}^t} = \frac{\partial \ell}{\partial \mathbf{h}^d} \underbrace{\frac{\partial \mathbf{h}^d}{\partial \mathbf{h}^{d-1}} \dots \frac{\partial \mathbf{h}^{t+1}}{\partial \mathbf{h}^t}}_{\text{Multiplication of } d-t \text{ matrices}} \frac{\partial \mathbf{h}^t}{\partial \mathbf{W}^t}$$

Multiplication of  $d-t$  matrices

# Two Issues for Deep Neural Networks

- Two common issues with  $\prod_{i=t}^{d-1} \frac{\partial \mathbf{h}^{i+1}}{\partial \mathbf{h}^i}$

Gradient Exploding



$$1.5^{100} \approx 4 \times 10^{17}$$

Gradient Vanishing



$$0.8^{100} \approx 2 \times 10^{-10}$$

# Example: FFN

---

- Assume FFN (without bias for simplicity)

$$f_t(\mathbf{h}^{t-1}) = \sigma(\mathbf{W}^t \mathbf{h}^{t-1}) \quad \sigma \text{ is the activation function}$$

$$\frac{\partial \mathbf{h}^t}{\partial \mathbf{h}^{t-1}} = \text{diag}(\sigma'(\mathbf{W}^t \mathbf{h}^{t-1})) (\mathbf{W}^t)^T \quad \sigma' \text{ is the gradient function of } \sigma$$

$$\prod_{i=t}^{d-1} \frac{\partial \mathbf{h}^{i+1}}{\partial \mathbf{h}^i} = \prod_{i=t}^{d-1} \text{diag}(\sigma'(\mathbf{W}^i \mathbf{h}^{i-1})) (\mathbf{W}^i)^T$$



# Gradient Exploding

- Use ReLU as the activation function

$$\sigma(x) = \max(0, x) \quad \text{and} \quad \sigma'(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

- Elements of  $\prod_{i=t}^{d-1} \frac{\partial \mathbf{h}^{i+1}}{\partial \mathbf{h}^i} = \prod_{i=t}^{d-1} \text{diag}(\sigma'(\mathbf{W}^i \mathbf{h}^{i-1})) (\mathbf{W}^i)^T$  may  
from  $\prod_{i=t}^{d-1} (\mathbf{W}^i)^T$

- Leads to large values when  $d-t$  is large

$$1.5^{100} \approx 4 \times 10^{17}$$

# Issues with Gradient Exploding


---

- Value out of range: infinity value
  - Severe for using 16-bit floating points
    - Range:  $6E-5 \sim 6E4$
- Sensitive to learning rate (LR)
  - Not small enough LR  $\rightarrow$  large weights  $\rightarrow$  larger gradients
  - Too small LR  $\rightarrow$  No progress
  - May need to change LR dramatically during training

# Gradient Vanishing

- Use sigmoid as the activation function

$$\sigma(x) = \frac{1}{1 + e^{-x}} \quad \sigma'(x) = \sigma(x)(1 - \sigma(x))$$



# Gradient Exploding

---

- Use sigmoid as the activation function

$$\sigma(x) = \frac{1}{1 + e^{-x}} \quad \sigma'(x) = \sigma(x)(1 - \sigma(x))$$

- Elements  $\prod_{i=t}^{d-1} \frac{\partial \mathbf{h}^{i+1}}{\partial \mathbf{h}^i} = \prod_{i=t}^{d-1} \text{diag}(\sigma'(\mathbf{W}^i \mathbf{h}^{i-1}))(\mathbf{W}^i)^T$  are products of  $d-t$  small values

$$0.8^{100} \approx 2 \times 10^{-10}$$

# Issues with Gradient Vanishing

---

- Gradients with value 0
  - Severe with 16-bit floating points
- No progress in training
  - No matter how to choose learning rate
- Severe with bottom layers
  - Only top layers are well trained
  - No benefit to make networks deeper

# Stabilize Training




# Stabilize Training

---

- Goal: make sure gradient values are in a proper range
  - E.g. in  $[1e-6, 1e3]$
- Multiplication  $\rightarrow$  plus
  - ResNet, LSTM (later lecture)
- Normalize
  - Gradient clipping
  - Batch Normalization / Layer Normalization (later)
- Proper weight initialization and activation functions

# Weight Initialization

- Initialize weights with random values in a proper range
- The beginning of training easily suffers to numerical instability
  - The surface far away from an optimal can be complex
  - Near optimal may be flatter
- Initializing according to  $\mathcal{N}(0, 0.01)$  works well for small networks, but not guarantee for deep neural networks





# Constant Variance for each Layer

---

- Treat both layer outputs and gradients are random variables
- Make the mean and variance for each layer's output are same, similar for gradients

Forward

$$\begin{aligned}\mathbb{E}[h_i^t] &= 0 \\ \text{Var}[h_i^t] &= a\end{aligned}$$

Backward

$$\mathbb{E}\left[\frac{\partial \ell}{\partial h_i^t}\right] = 0 \quad \text{Var}\left[\frac{\partial \ell}{\partial h_i^t}\right] = b \quad \forall i, t$$

$a$  and  $b$  are constants

# Example: FFN

---

- Assumptions  $\mathbb{E}[w_{i,j}^t] = 0$ ,  $\text{Var}[w_{i,j}^t] = \gamma_t$ 
  - i.i.d  $w_{i,j}^t$
  - $h_i^{t-1}$  is independent to  $w_{i,j}^t$
  - identity activation:  $\mathbf{h}^t = \mathbf{W}^t \mathbf{h}^{t-1}$  with  $\mathbf{W}^t \in \mathbb{R}^{n_t \times n_{t-1}}$

$$\mathbb{E}[h_i^t] = \mathbb{E} \left[ \sum_j w_{i,j}^t h_j^{t-1} \right] = \sum_j \mathbb{E}[w_{i,j}^t] \mathbb{E}[h_j^{t-1}] = 0$$

# Forward Variance

---

$$\begin{aligned}\text{Var}[h_i^t] &= \mathbb{E}[(h_i^t)^2] - \mathbb{E}[h_i^t]^2 = \mathbb{E} \left[ \left( \sum_j w_{i,j}^t h_j^{t-1} \right)^2 \right] \\ &= \mathbb{E} \left[ \sum_j \left( w_{i,j}^t \right)^2 \left( h_j^{t-1} \right)^2 + \sum_{j \neq k} w_{i,j}^t w_{i,k}^t h_j^{t-1} h_k^{t-1} \right] \\ &= \sum_j \mathbb{E} \left[ \left( w_{i,j}^t \right)^2 \right] \mathbb{E} \left[ \left( h_j^{t-1} \right)^2 \right] \\ &= \sum_j \text{Var}[w_{i,j}^t] \text{Var}[h_j^{t-1}] = n_{t-1} \gamma_t \text{Var}[h_j^{t-1}] \quad \Rightarrow \quad n_{t-1} \gamma_t = 1\end{aligned}$$

$n_{t-1}$  is the number of units in t-1 layer

# Backward Mean and Variance

---

- Apply forward analysis as well

$$\frac{\partial \ell}{\partial \mathbf{h}^{t-1}} = \frac{\partial \ell}{\partial \mathbf{h}^t} \mathbf{W}^t \quad \text{leads to} \quad \left( \frac{\partial \ell}{\partial \mathbf{h}^{t-1}} \right)^T = (\mathbf{W}^t)^T \left( \frac{\partial \ell}{\partial \mathbf{h}^t} \right)^T$$

$$\mathbb{E} \left[ \frac{\partial \ell}{\partial h_i^{t-1}} \right] = 0$$

$$\text{Var} \left[ \frac{\partial \ell}{\partial h_i^{t-1}} \right] = n_t \gamma_t \text{Var} \left[ \frac{\partial \ell}{\partial h_j^t} \right] \quad \Rightarrow \quad n_t \gamma_t = 1$$


# Xavier Initialization

---


- Conflict goal to satisfies both  $n_{t-1}\gamma_t = 1$  and  $n_t\gamma_t = 1$
- Xavier  $\gamma_t(n_{t-1} + n_t)/2 = 1 \rightarrow \gamma_t = 2/(n_{t-1} + n_t)$ 
  - Normal distribution  $\mathcal{N}\left(0, \sqrt{2/(n_{t-1} + n_t)}\right)$
  - Uniform distribution  $\mathcal{U}\left(-\sqrt{6/(n_{t-1} + n_t)}, \sqrt{6/(n_{t-1} + n_t)}\right)$ 
    - Variance of  $\mathcal{U}[-a, a]$  is  $a^2/3$
- Adaptive to weight shape, especially when  $n_t$  varies

# Other heuristics: Early stopping

- Continued training can result in over fitting to training data
  - Track performance on a held-out validation set
  - Apply one of several early-stopping criterion to terminate training when performance on validation set degrades significantly



# Additional heuristics: Gradient clipping



- Often the derivative will be too high
  - When the divergence has a steep slope
  - This can result in instability
- **Gradient clipping**: set a ceiling on derivative value

$$\text{if } \partial_w D > \theta \text{ then } \partial_w D = \theta$$

- Typical  $\theta$  value is 5
- Can be easily set in pytorch/tensorflow

# Recap

---

- Numerical issues in training
  - gradient explosion
  - gradient vanishing
- Proper initialization of parameters



# Next Up

---

- Convolutional Neural Networks
- Visual perception:
  - Image classification
  - Object recognition
  - Face detection